

Kaushik Subramanian

CONTACT INFORMATION	1270 West Peachtree Street NW, 21E Arts Center Tower, Atlanta, GA - 30309, USA	Mobile: +1 650 454 4701 E-mail: ksubrama@cc.gatech.edu www: http://www.cc.gatech.edu/~ksubrama
RESEARCH OBJECTIVES	To perform research in the field of Interactive Machine Learning with the goal of designing learning algorithms that can efficiently interact with humans and utilize their information to solve real-world problems.	
ACADEMIC QUALIFICATIONS	Georgia Institute of Technology, USA Doctor of Philosophy in Robotics (Computer Science)	August 2010 - Present
	Rutgers, The State University of New Jersey, USA Master of Science in Electrical and Computer Engineering	August 2008 - May 2010
	Anna University, India Bachelor of Engineering in Electrical and Communication Engineering	July 2004 - May 2008
RESEARCH EXPERIENCE	<i>Intern</i> at Microsoft Research, Redmond, USA Completed a 3 month internship on applying machine learning algorithms on data from human muscle activity.	June 2011 - August 2011
	<i>Research Assistant</i> in the Socially Intelligent Machines Lab (SIM), GAtech.	August 2010 - Present
	<i>Research Assistant</i> in the pFunk Machine Learning Lab, GAtech.	August 2010 - Present
	<i>Adjunct Member</i> of the Real Life Reinforcement Learning Lab (RL ³), Rutgers.	January 2009 - May 2010
	<i>Intern</i> at RWTH Aachen University, Germany Completed a 3 month internship on humanoid robot learning by demonstration using gaussian mixture models.	June 2009 - August 2009
	<i>Intern</i> at Texas Instruments, India Completed a 3 month internship focusing on data compression in video encoding techniques.	May 2008 - July 2008
	<i>Research Trainee</i> at WArAn Research FoundaTion, India Completed a 2 year research training program with specialization in signal processing.	June 2006 - May 2008
MASTERS DISSERTATION	“ <i>HELP - Human assisted Efficient Learning Protocols</i> ” - an analysis into the effect of human interactions on learning algorithms, with focus on methods like learning by demonstration and apprenticeship learning. (Abstract attached at the end) <i>Advisor - Prof. Michael Littman</i>	
PUBLICATIONS	Monica Babes, Vukosi Marivate, Kaushik Subramanian, Michael L. Littman: <i>Apprenticeship Learning about Multiple Intentions</i> . Appeared in ICML 2011, Bellevue, Washington, USA, June 2011.	
	Thomas J. Walsh, Kaushik Subramanian, Michael L. Littman, Carlos Diuk: <i>Generalizing Apprenticeship Learning across Hypothesis Classes</i> . Appeared in ICML 2010, Haifa, Israel, June 2010.	
WORKSHOPS	Kaushik Subramanian, Andrea Thomaz, Charles Isbell: <i>Learning Options through Human Interaction</i> . Appeared in the ALIHT Workshop, IJCAI 2011, Barcelona, Spain, July 2011.	
	Kaushik Subramanian: <i>Task Space Behavior Learning for Humanoid Robots using Gaussian Mixture Regression</i> . Appeared in AAAI 2010, Atlanta, USA, July 2010.	
TECHNICAL REPORTS	Baris Akgun, Kaushik Subramanian, Jaeun Shim, Andrea Thomaz : <i>Learning Tasks and Skills Together From a Human Teacher</i> . Appeared in AAAI 2011, San Francisco, USA, August 7th 2011.	
	Kaushik Subramanian and Michael Littman: <i>Efficient Apprenticeship Learning with Smart Humans</i> . Appeared in AAAI 2010, Atlanta, USA, July 2010.	
PROJECTS IN PROGRESS	<i>Feature Construction using Human Interaction</i> (GAtech) An exploration into the existing methods used to construct features from low-level data using human interaction. The purpose is to automatically learn features and task-specific abstractions that assist learning	September 2011

algorithms and additionally reduce the burden on the human-designer. *Advisor - Prof. Charles Isbell and Prof. Andrea Thomaz*

Novel Interaction Strategies for Learning from Teleoperation (GAtch) September 2011
(Under Review at ICRA 2011) Aims to overcome the challenges associated with using teleoperation as an input modality for Robot learning from demonstration. A novel demonstration strategy is proposed using the concept of keyframes where demonstrations are in the form of a discrete set of robot configurations. Studies show that keyframes are intuitive to the users and are particularly useful in providing noise-free demonstrations when demonstrating tasks to a robot by teleoperation. *Advisor - Prof. Andrea Thomaz*

Using Human Task Decompositions to enhance Monte Carlo Tree Search (GAtch) May 2011
Developed an automatic way of leveraging human input to acquire task decompositions. These can be used as pruning heuristics and rollout policies to assist Monte Carlo Tree Search. They were applied to two game-based domains - PacMan and Mouse. The enhanced MCTS approach outperforms standard approaches in terms of optimality of the solution by 35%. *Advisor - Prof. Charles Isbell and Prof. Andrea Thomaz*

COMPLETED
PROJECTS

Robot Learning: Kinesthetic Teaching vs. Teleoperation (GAtch) April 2011
Two interaction methods were compared, kinesthetic teaching and teleoperation, by allowing users to show demonstrations of different skills. The results were evaluated using skill dependent quantitative measures, timing information and survey questions. The results show that kinesthetic teaching is the faster, easier and more accurate input modality and highlights the areas where teleoperation can be improved. *Advisor - Prof. Andrea Thomaz*

Learning Attributes from Human Interaction (GAtch) March 2011
Used supervised learning algorithms like kNN and GMMs, to learn relevant attributes for a simulated animal classification and a robot table clean-up task. The algorithm utilized human interaction to instantiate the attributes that aided complex concept learning and increased overall task efficiency. *Advisor - Prof. Andrea Thomaz*

MDP-based Planning for a Table-top Search and Find Task (GAtch) December 2010
A novel tree-based task representation was developed to perform table-top search of occluded objects. The problem was modeled as a POMDP and solved to acquire the set of actions that lead to successful object retrieval. *Advisor - Prof. Mike Stilman*

AAAI 2010 Learning by Demonstration Challenge (GAtch) July 2010
A live demonstration was performed of the "Taxi Task" using Apprenticeship Learning. A MindstormsTM Lego robot was navigated through a discrete world and it was allowed to build a model of the environment. The robot was able to learn the optimal policy from a single demonstration. *Advisor - Prof. Michael Littman*

Interactive Learning with the Highway Car Domain (RL³ Lab) December 2009
A novel approach was developed where the humans provides high-level state abstractions to learn the task of navigating on a simulated highway. The criteria used by the human was - "states are similar if the same optimal action is to performed in both the states". The interactive abstraction significantly sped-up the performance of the agent. *Advisor - Prof. Michael Littman*

Robot Learning by Demonstration using GMM's (KBSG Lab, RWTH University) July 2009
A behavior acquisition model was developed for the Nao's using Gaussian Regression. After generalizing the kinesthetic demonstrations, the robot was used to imitate constrained reaching gestures. *Advisor - Prof. Gerhard Lakemeyer*

Best Narration Award - Introduction to Reinforcement Learning (RL³ Lab) April 2009
The Lego Mindstorms was programmed to learn real-time in a deterministic environment and to build a model of the world using concepts of Graph Search and Dynamic Programming. A video tutorial was submitted to IJCAI 2009. *Advisor - Prof. Michael Littman*

Autonomous Object Recognition using Corner Detection (Rutgers) December 2008
Implemented using Corner Descriptors and Geometric Point Matching methods. The advantage of the system was the reduced number of descriptor points as compared to the SIFT algorithm. *Advisor - Prof. Lawrence Rabiner*

COMPUTING SKILLS

Programming - C, C++, Java, Python, Matlab
Softwares - ROS, OpenCV, Microsoft Robotics Developer Studio, Tekkotsu, Fawkes
Assembly Language - AVR Microcontroller, 8051, 8086, 8085
Operating Systems - Unix and Windows